




Cloudera Improvements in Apache Spark

Brian Baillod | Sales Engineer



Agenda

- Introduction
- Spark One Platform Initiative
- Spark Overview and Improvements
- Spark Proof of Concept
- Kudu and Record Service

Cloudera company snapshot

Founded 2008, by former employees of ORACLE YAHOO! facebook Google

Employees Today 900+ worldwide

World Class Support More than 75 24x7 global staff

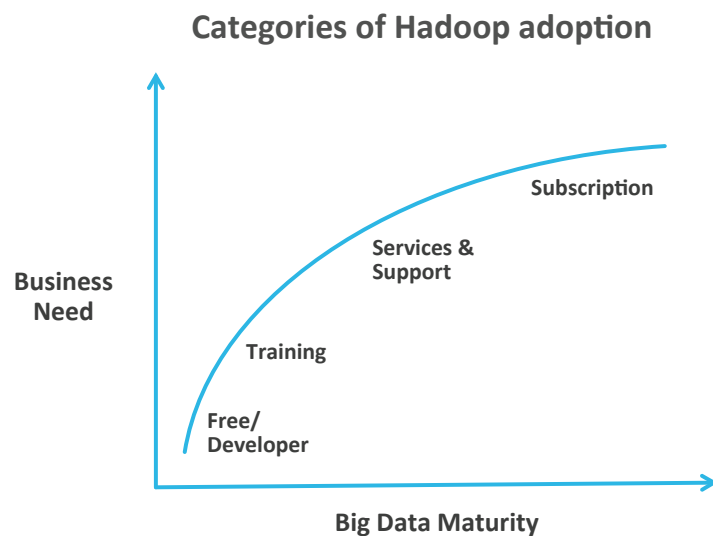
Cloudera University Over 40,000 trained

We help code Hadoop Cloudera employees are leading developers & contributors to the complete Apache Hadoop ecosystem of projects

We help fix Hadoop Cloudera fixed 60% of all Hadoop JIRA bugs

cloudera

Hadoop Adoption



Free/Developer ➡ Over **2.5 million** downloads

Training ➡ **60% of Fortune 100** attended Cloudera training, **over 40,000** trained since 2009

Service & Support ➡ **9/10** for support satisfaction, ability to solve technical issues **#1 recommendation**

Subscription ➡ Over **2x** revenue of nearest competitor, **90%** renewal rate

What is Spark


- Fast general purpose processing engine for large data
- Provides API's in Java, Scala and Python
- Includes an advanced DAG execution engine that supports in-memory computing
- Includes high level tools like SparkSQL, Mllib, GraphX, and Spark Streaming
- Can run in a cluster, standalone, or local
- Latest version is 1.5.1
- Spark.apache.org
- LOTS of momentum

Cloudera One Platform Initiative

- Cloudera is doubling down on Spark
- Outlining a vision for the future
 - Kudu, Record Service, Auto-tuning, Security, Kafka integration
- Challenging other vendors to participate in Spark Development

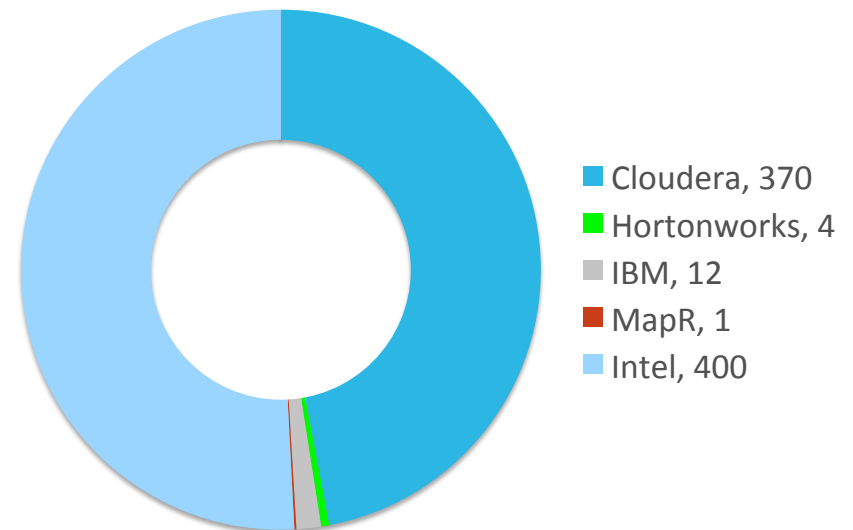
Cloudera's Engineering Commitment to Spark

Spark Committers by Hadoop Distribution*



* IBM and MapR have **0** committers

Spark Patches by Hadoop Distribution

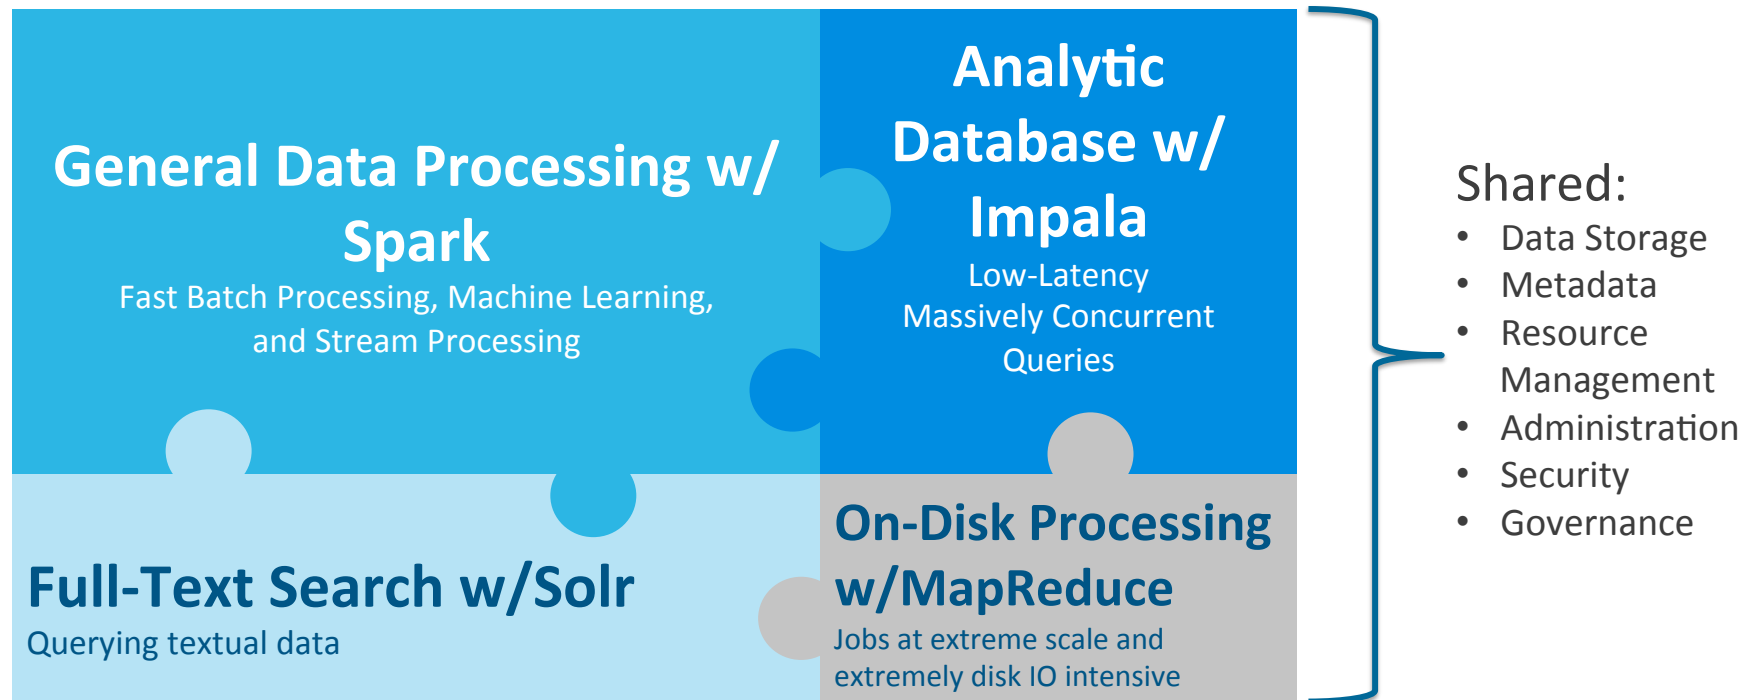


Spark will replace MapReduce

To become the standard execution engine for Hadoop

The Future of Data Processing on Hadoop

Spark complemented by specialized fit-for-purpose engines




Why is Cloudera leading this initiative?

- Cloudera was the **first** Hadoop vendor to ship and support Spark
- Spark is a **fully integrated** part of Cloudera's platform
 - Shared data, metadata, resource management, administration, security, and governance
- Cloudera is the **first** Hadoop vendor to offer Spark training
 - Trained more customers than any other vendor
- Cloudera has **more Spark customers** in production than all other companies combined



Spark Overview and Improvements



Apache Spark

Flexible, in-memory data processing for Hadoop

Easy Development

- Rich APIs for Scala, Java, and Python
- Interactive shell

Flexible Extensible API

- APIs for different types of workloads:
 - Batch
 - Streaming
 - Machine Learning
 - Graph

Fast Batch & Stream Processing

- In-Memory processing and caching

Easy Development

High Productivity Language Support

Python

```
lines = sc.textFile(...)
lines.filter(lambda s: "ERROR" in s).count()
```

Scala

```
val lines = sc.textFile(...)
lines.filter(s => s.contains("ERROR")).count()
```

Java

```
JavaRDD<String> lines = sc.textFile(...);
lines.filter(new Function<String, Boolean>() {
    Boolean call(String s) {
        return s.contains("error");
    }
}).count();
```

- Native support for multiple languages with identical APIs
 - Scala, Java, Python
- Use of closures, iterations, and other common language constructs to minimize code
 - 2-5x less code

Python Or Scala?

- Use Python for prototyping
- Spark Python API is slower than Scala
- Use Scala for development
 - Steep learning curve for functional programming

Easy Development

Use Interactively

[illegible]

- Interactive exploration of data for data scientists
 - No need to develop “applications”
- Developers can prototype application on live system


Easy Development

Expressive API

- map
- filter
- groupBy
- sort
- union
- join
- leftOuterJoin
- rightOuterJoin
- reduce
- count
- fold
- reduceByKey
- groupByKey
- cogroup
- cross
- zip
- sample
- take
- first
- partitionBy
- mapWith
- pipe
- save
- ...

Memory Management for Greater Performance

Memory can be enabler for high performance big data applications



Trends:

- ½ price every 18 months
- 2x bandwidth every 3 years

Spark Concepts

- RDD – Resilient Distributed Dataset
- Transformations
- Actions
- Caching
- DataFrames
- Spark Streaming
- SparkSQL
- Pluggable Spark



Resilient Distributed Dataset (RDD)

- Read-only partitioned collection of records
- Created through:
 - Transformation of data in storage
 - Transformation of RDDs
- Contains lineage to compute from storage
- Lazy materialization
- Users control persistence and partitioning

RDD Operations

- **Transformations** create new RDD from an existing one
- **Actions** run computation on RDD and return a value
- Transformations are lazy
- Actions materialize RDDs by computing transformations
- RDDs can be cached to avoid re-computing

Example Operations

Transformations

- Map
- Filter
- Sample
- Join

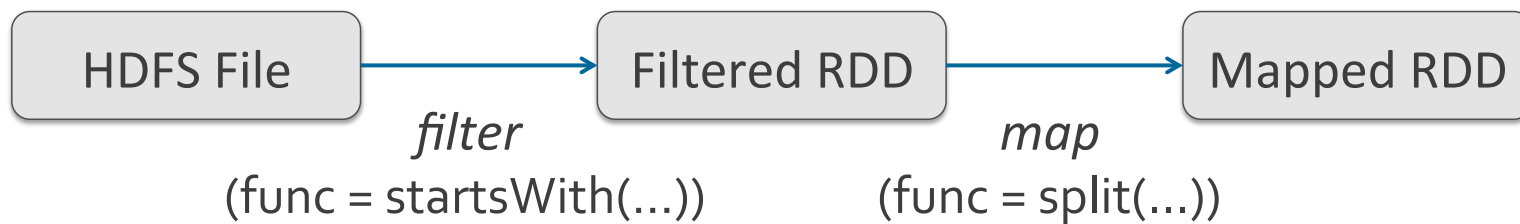
Actions

- Reduce
- Count
- First, Take
- SaveAs

Fault-Tolerance

- RDDs contain lineage
 - Lineage: Source location and list of transformations
 - Lost partitions can be re-computed from source data

```
msgs = textFile.filter(lambda s: s.startsWith("ERROR"))  
                .map(lambda s: s.split("\t")[2])
```



Caching – Storage Levels

Different options provide tradeoffs between memory usage and CPU efficiency. Cache when using iterative algorithms.

- MEMORY_ONLY – most CPU efficient, data has to fit in memory
- MEMORY_ONLY_SER – More space efficient but still reasonably fast
- MEMORY_AND_DISK
- MEMORY_AND_DISK_SER
- DISK_ONLY
- MEMORY_ONLY_2, MEMORY_AND_DISK_2...

Data Frames

- Distributed collection of rows organized into named columns
- Spark SQL's Data Source API can read and write Data Frames using a variety of formats
 - Hive, JSON, Parquet, HDFS
- Calling the DataFrame API can let you
 - Select the columns you want
 - Join data sources
 - Aggregate and Filter
- Spark 1.5 lets you access the Hive Metastore to read/write schemas directly.

Spark Streaming

What is it?

- Run *continuous* processing of data using Spark's core API
- Extends Spark concepts to fault-tolerant, transformable streams
- Adds "rolling window" operations
 - Example: Compute rolling averages or counts for data over last five minutes

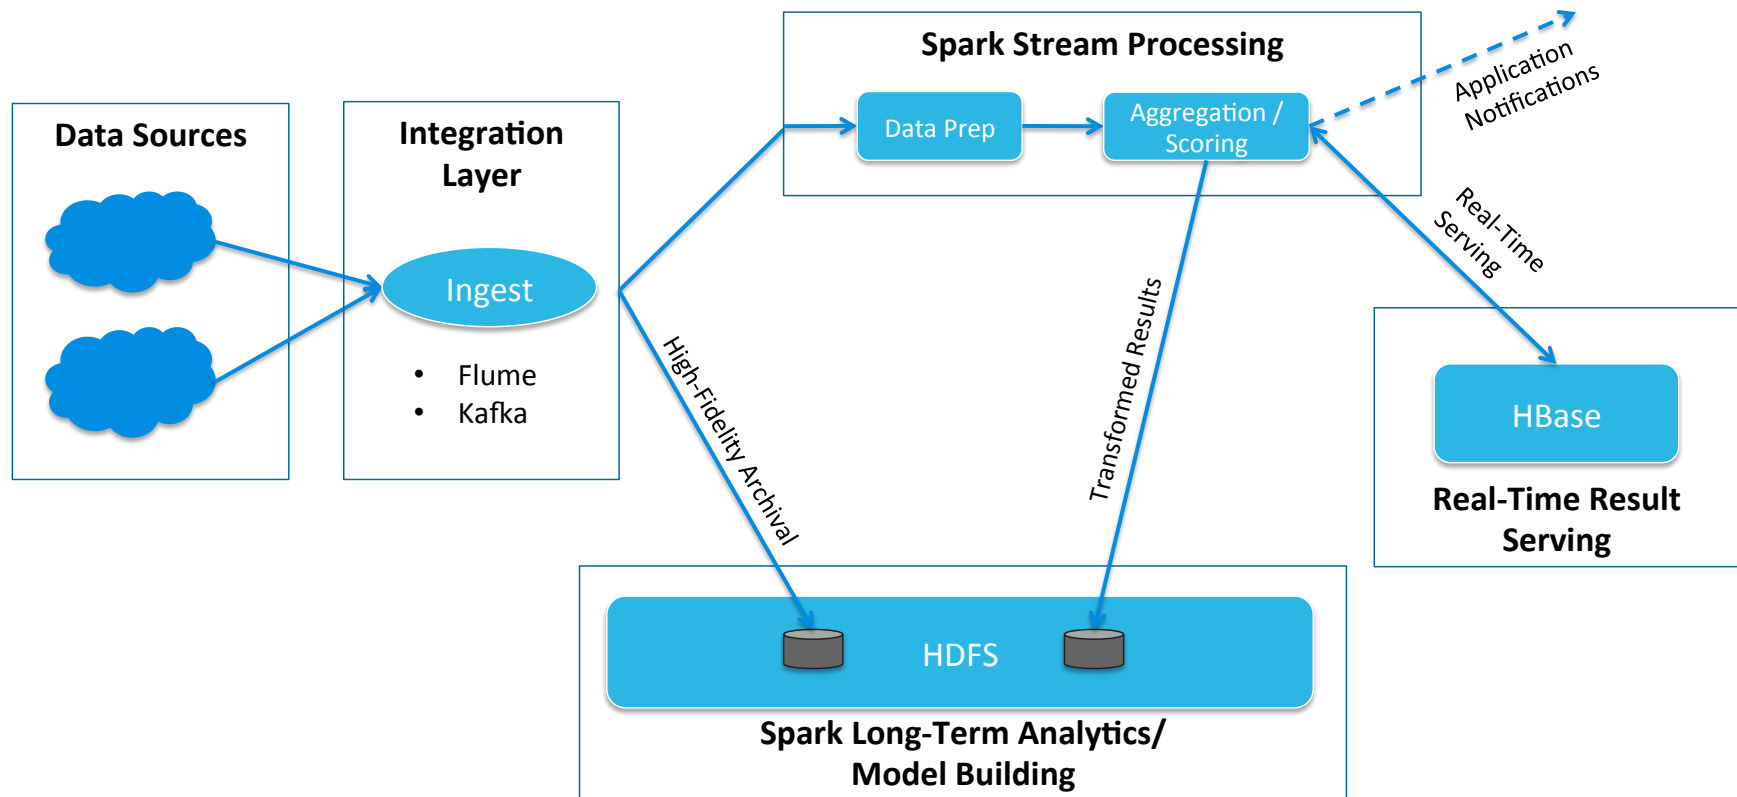
Common Use Cases:

- "On-the-fly" ETL as data is ingested into Hadoop/HDFS
- Detect anomalous behavior and trigger alerts
- Continuous reporting of summary metrics for incoming data

Benefits:

- Same programming paradigm for streaming and batch
- Excellent throughput
 - Scale easily to support large volumes of data ingest

Spark Streaming Architectures




SparkSQL

Machine Learning Applications

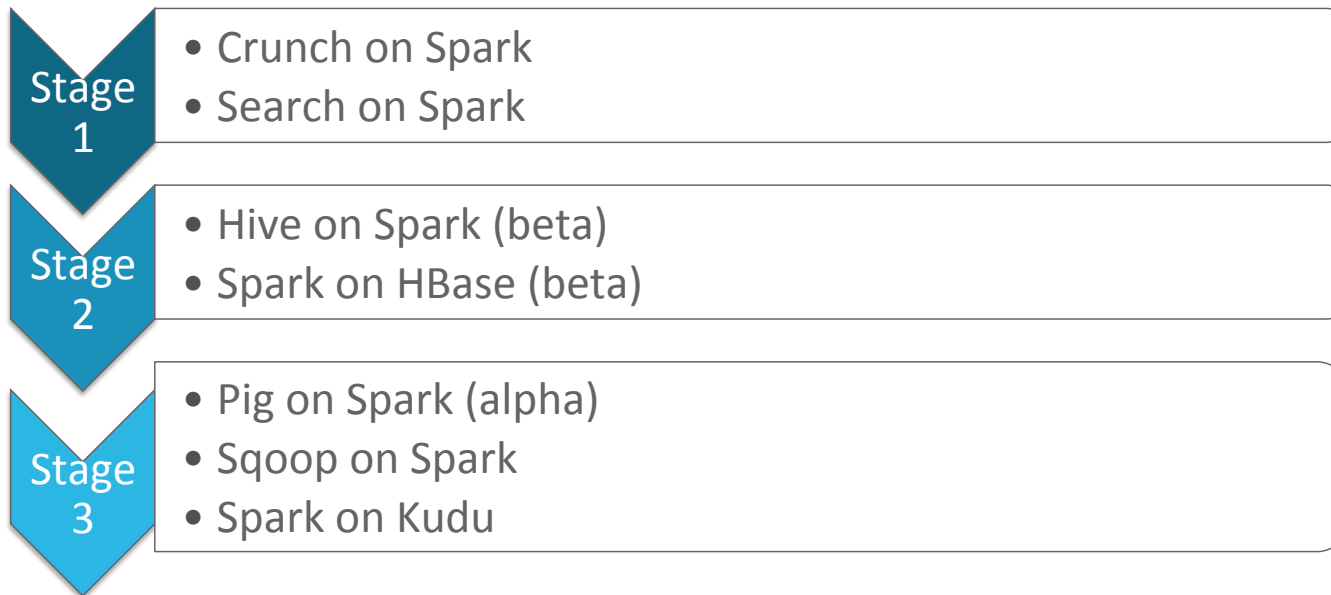
- Goal:
 - Spark/Java Developers and Data Scientists can inline SQL into Spark apps
- Designed for:
 - Ease of development for Spark developers
 - Handful of concurrent Spark jobs
- Strengths:
 - Ease of embedding SQL into Java or Scala applications
 - SQL for common functionality in developer flow (eg. aggregations, filters, samples)

Impala Remains Tool of Choice for Interactive SQL



Pluggable Spark – replace MapReduce

Cloudera is leading community development to port components to Spark:



Spark Customer Use Cases

Core Spark



Financial
Services

- Portfolio Risk Analysis
- ETL Pipeline Speed-Up
- 20+ years of stock data



Health

- Identify disease-causing genes in the full human genome
- Calculate Jaccard scores on health care data sets



ERP

- Optical Character Recognition and Bill Classification



Data
Services

- Trend analysis
- Document classification (LDA)
- Fraud analytics

cloudera

Spark Streaming



Financial
Services

- Online Fraud Detection



Health

- Incident Prediction for Sepsis



Retail

- Online Recommendation Systems
- Real-Time Inventory Management

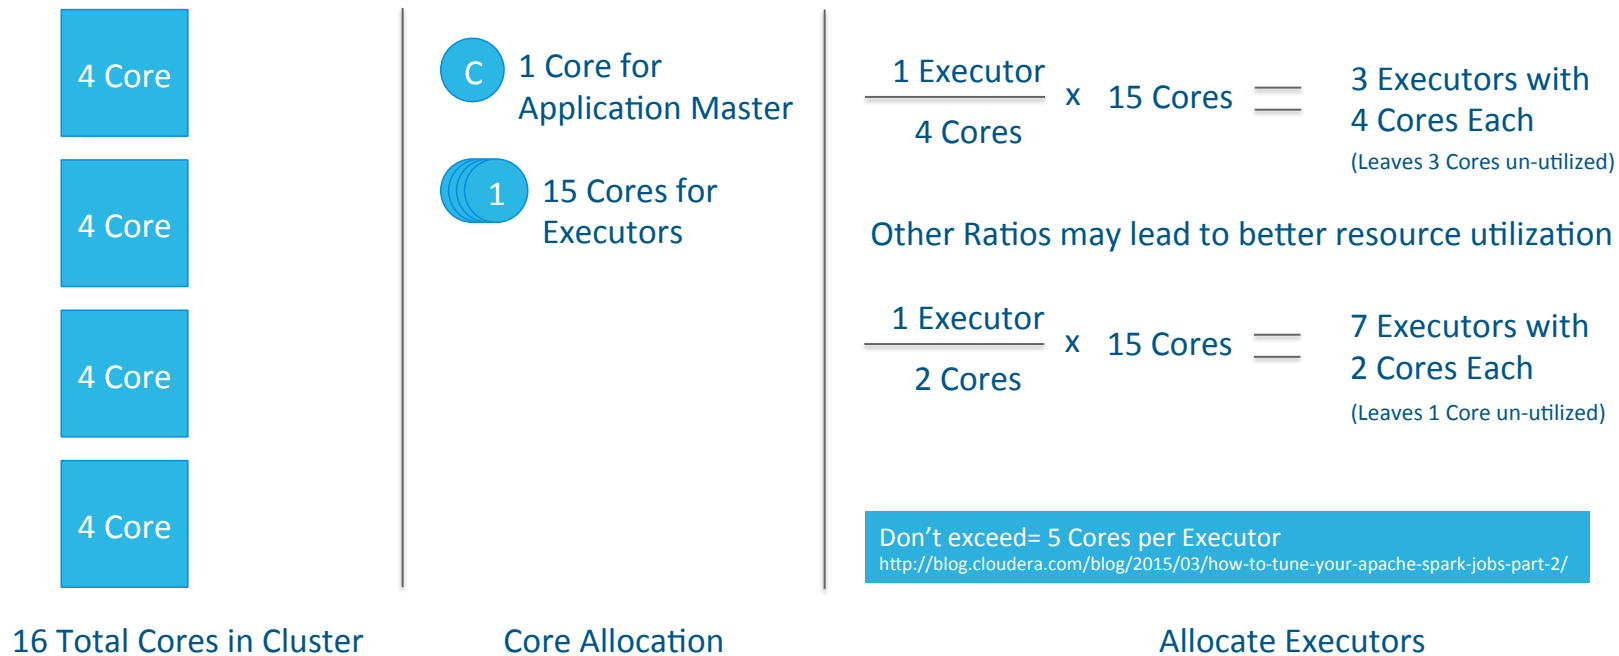


Ad Tech

- Real-Time Ad Performance Analysis


Doing the Math – Executors and Cores

Determine the optimal resource allocation for the Spark job

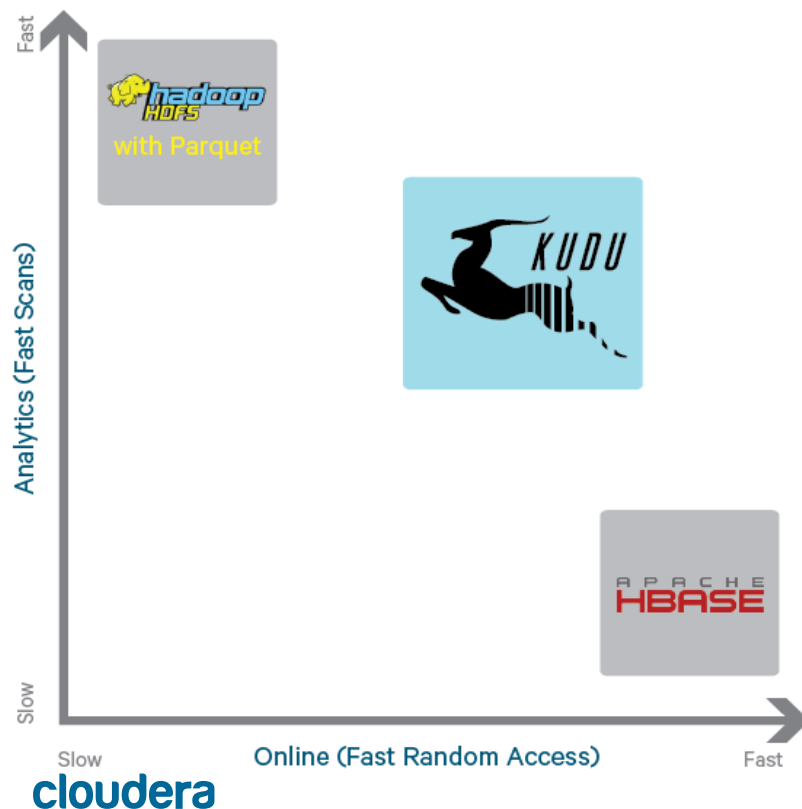


cloudera®

Kudu



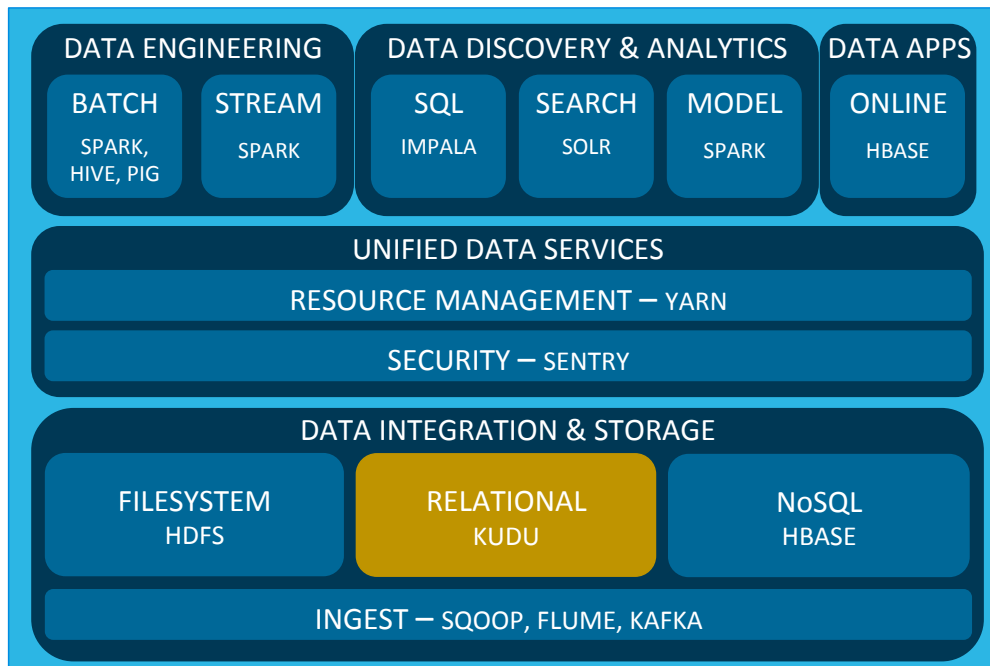
Kudu Design Goals



- **High throughput** for big scans (columnar storage and replication)
Goal: Within 2x of Parquet
- **Low-latency** for short accesses (primary key indexes and quorum design)
Goal: 1ms read/write on SSD
- **Database-like** semantics (initially single-row ACID)
- **Relational data model**
 - SQL query
 - “NoSQL” style scan/insert/update (Java client)

Kudu

Storage for Fast Analytics on Fast Data



- New updating column store for Hadoop
 - Simplifies the architecture for building analytic applications on changing data
 - Designed for fast analytic performance
 - Natively integrated with Hadoop
- Apache-licensed open source (intent to donate to ASF)
- Beta now available

Kudu Trade-Offs

- Random updates will be slower
 - HBase model allows random updates without incurring a disk seek
 - Kudu requires a key lookup before update, Bloom lookup before insert
- Single-row reads may be slower
 - Columnar design is optimized for scans
 - **Future:** may introduce “column groups” for applications where single-row access is more important

Resources

Join the community

<http://getkudu.io>

Download the Beta

cloudera.com/downloads


Read the Whitepaper

getkudu.io/kudu.pdf



cloudera®

RecordService

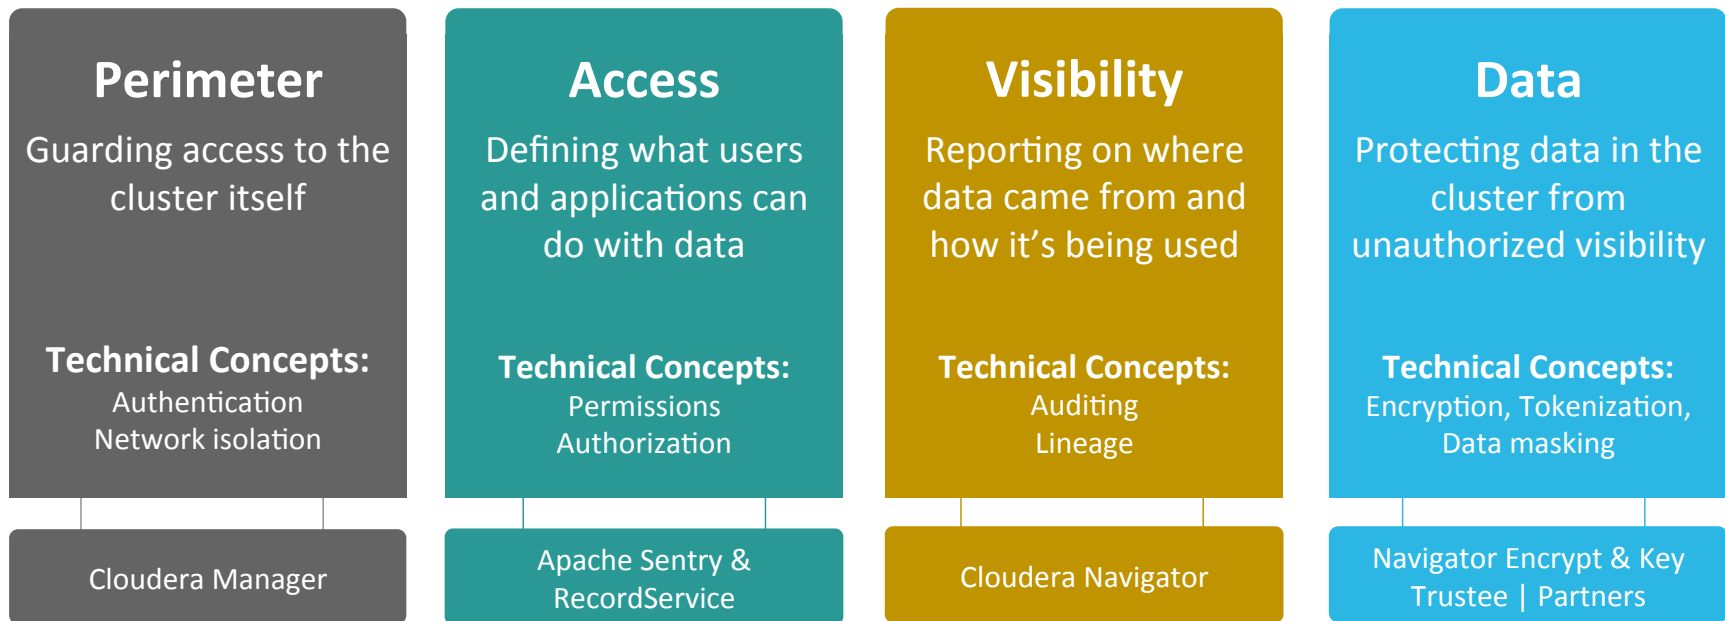


Hadoop started out with zero security

- Didn't need it for the Silicon Valley applications
- Does need it for Corporate applications
- Cloudera is working on providing full featured Spark Security

Comprehensive, Compliance-Ready Security

Authentication, Authorization, Audit, and Compliance



Active Directory and Kerberos

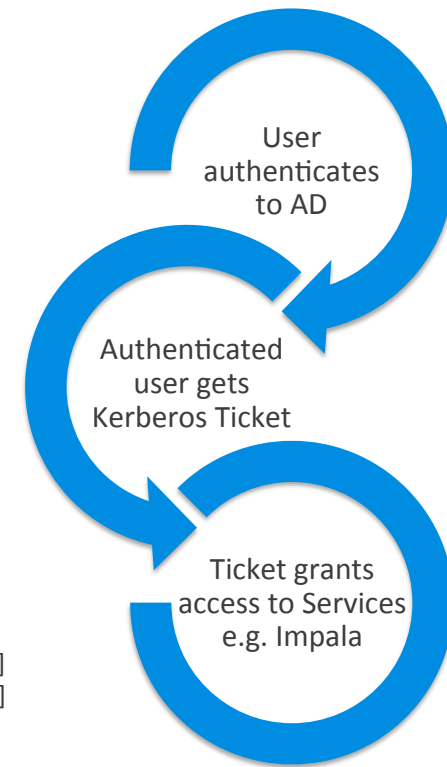
Active Directory

- Manages Users, Groups, and Services
- Provides username / password authentication
- Group membership determines Service access

Kerberos

- Trusted and standard third-party
- Authenticated users receive “Tickets”
- “Tickets” gain access to Services


User [ssmith]
Password[*****]



Fine-Grained Access Control in HDFS

Across All Hadoop Paths

Columns:

Sensitive column visibility varies by role (Ex. credit card numbers)

- Managers: 1234 5678 1234 5678
- Call Center: XXXX XXXX XXXX 5678
- Analysts: XXXX XXXX XXXX XXXX
- Others: No access to credit card column

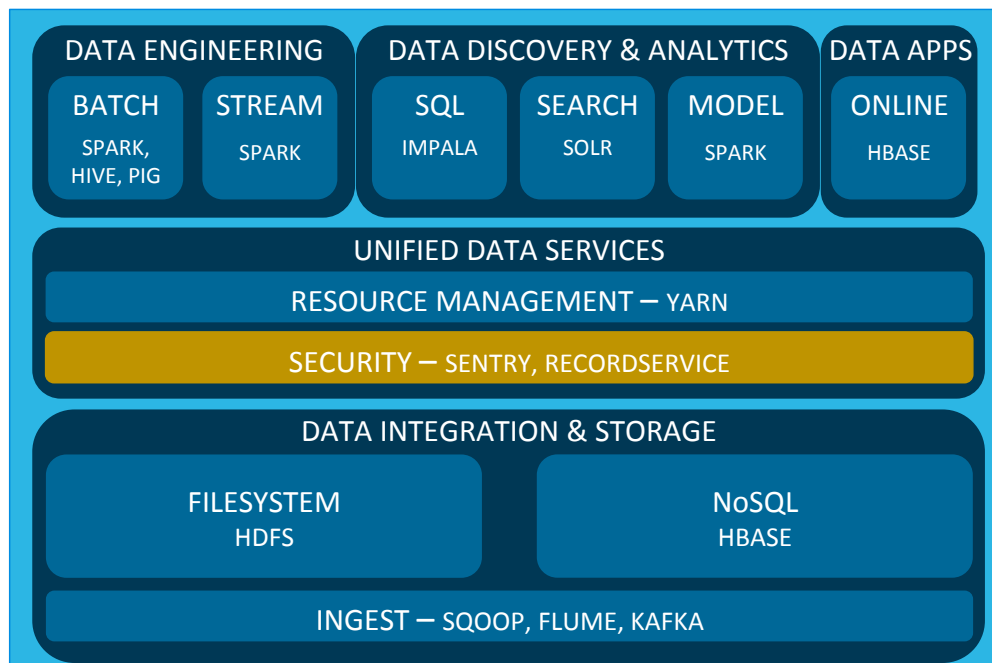
Rows:

Different user groups need access to different records

- European privacy laws
- Government security clearance
- Financial information restrictions

RecordService

Unified Access Control Enforcement



- New high performance security layer that centrally enforces access control policies across Hadoop
 - Complements Apache Sentry's unified policy definition
 - Row- and column-based security
 - Dynamic data masking
- Apache-licensed open source
- Beta now available

Fine-Grained HDFS Access without RecordService

Split the original file
Use HDFS permissions to limit access

Date/time	Accnt #	SSN	Asset	Trade	Country
09:33:11 16-Feb-2015	0234837823	238-23-9876	AAPL	Sell	US
11:33:01 16-Feb-2015	3947848494	329-44-9847	TBT	Buy	EU
14:12:34 16-Feb-2015	4848367383	123-56-2345	IBM	Sell	UK
09:22:03 16-Feb-2015	3485739384	585-11-2345	INTC	Buy	US
11:55:33 16-Feb-2015	3847598390	234-11-8765	F	Buy	US
10:22:55 16-Feb-2015	8765432176	344-22-9876	UA	Buy	UK
13:45:24 16-Feb-2015	3456789012	412-22-8765	AMZN	Sell	EU
09:03:44 16-Feb-2015	4857389329	123-44-5678	TMV	Buy	US
15:55:55 16-Feb-2015	4756983234	234-76-9274	MA	Buy	UK

Date/time	Accnt #	SSN	Asset	Trade	Country
09:33:11 16-Feb-2015	0234837823	238-23-9876	AAPL	Sell	US
09:22:03 16-Feb-2015	3485739384	585-11-2345	INTC	Buy	US
11:55:33 16-Feb-2015	3847598390	234-11-8765	F	Buy	US
09:03:44 16-Feb-2015	4857389329	123-44-5678	TMV	Buy	US

Date/time	Accnt #	SSN	Asset	Trade	Country
11:33:01 16-Feb-2015	3947848494	329-44-9847	TBT	Buy	EU
13:45:24 16-Feb-2015	3456789012	412-22-8765	AMZN	Sell	EU

Date/time	Accnt #	SSN	Asset	Trade	Country
14:12:34 16-Feb-2015	4848367383	123-56-2345	IBM	Sell	UK
10:22:55 16-Feb-2015	8765432176	344-22-9876	UA	Buy	UK
15:55:55 16-Feb-2015	4756983234	234-76-9274	MA	Buy	UK

Fine-Grained HDFS Access Control with RecordService

- Apply controls to the master data file
- Row, column, and sub-column (masking) controls
- Enforce these across all access paths

Column-Level Controls					
Date/time	Accnt #	SSN	Asset	Trade	Country
09:33:11 16-Feb-2015	0234837823	238-23-9876	AAPL	Sell	US
11:33:01 16-Feb-2015	3947848494	329-44-9847	TBT	Buy	EU
14:12:34 16-Feb-2015	4848367383	123-56-2345	IBM	Sell	EU
09:22:03 16-Feb-2015	3485739384	585-11-2345	INTC	Buy	US
11:55:33 16-Feb-2015	3847598390	234-11-8765	F	Buy	US
10:22:55 16-Feb-2015	8765432176	344-22-9876	UA	Buy	EU
13:45:24 16-Feb-2015	3456789012	412-22-8765	AMZN	Sell	EU

Row-Level Controls



What U.S. Brokers See

Column-Level Controls					
Date/time	Accnt #	SSN	Asset	Trade	Country
09:33:11 16-Feb-2015	0234837823	XXX-XX 9876	AAPL	Sell	US
09:22:03 16-Feb-2015	3485739384	XXX-XX 2345	INTC	Buy	US
11:55:33 16-Feb-2015	3847598390	XXX-XX 8765	F	Buy	US

Row-Level Controls

Spark Resources

- Learn Spark
 - [Spark Cookbook](#) – by Rishi Yadav
 - O'Reilly [Advanced Analytics with Spark](#) eBook (written by Clouderans)
 - [Cloudera Developer Blog](#)
 - cloudera.com/spark
- Get Trained
 - [Cloudera Spark Training](#)

cloudera

